

1 Latex Template for CO600 and CO620 Project
2 Reports

3 Example Name
em123@kent.ac.uk



School of Computing
University of Kent

Word Count: 6,100

4 June 4, 2020

6 This is a L^AT_EXtemplate for CO600/CO620 projects.

1 Introduction

Note the difference between `\cite` and `\citep` in this Latex template. `\citep` will give you (Simmons, 1988) whereas `\cite` will give Simmons (1988). The difference is important and is explained here <https://www.kent.ac.uk/learning/resources/studyguides/harvardreferencingquickguide.pdf>.

The paragraph that is above is important. Please read it!

There is some example text below that shows how to use Latex, and there are a few examples of `\cite` and `\citep`.

One of the most promising general approaches for solving combinatorial search problems is to generate an initial, suboptimal solution and then to apply local *repair* heuristics. Techniques based on this approach have met with empirical success on many combinatorial problems, including the traveling salesman and graph partitioning problems (Johnson, Papadimitrou and Yannakakis, 1988). Such techniques also have a long tradition in AI, most notably in problem-solving systems that operate by debugging initial solutions (Simmons, 1988; Sussman, 1975). In this paper, we describe how this idea can be extended to constraint satisfaction problems (CSPs) in a natural manner.

This box is an example todo note. Numbered lines are useful for marking. To remove line numbers, comment out the `lineno` package at the top of this latex file.

Most of the previous work on CSP algorithms has assumed a “constructive” backtracking approach in which a partial assignment to the variables is incrementally extended. In contrast, our method (Minton et al., 1990) creates a complete, but inconsistent assignment and then repairs constraint violations until a consistent assignment is achieved. The method is guided by a simple ordering heuristic for repairing constraint violations: identify a variable that is currently in conflict and select a new value that minimizes the number of outstanding constraint violations.

We present empirical evidence showing that on some standard problems our approach is considerably more efficient than traditional constructive backtracking methods. For example, on the n -queens problem, our method quickly finds solutions to the one million queens problem. We argue that

the reason that repair-based methods can outperform constructive methods is because a complete assignment can be more informative in guiding search than a partial assignment. However, the utility of the extra information is domain dependent. To help clarify the nature of this potential advantage, we present a theoretical analysis that describes how various problem characteristics may affect the performance of the method. This analysis shows, for example, how the “distance” between the current assignment and solution (in terms of the minimum number of repairs that are required) affects the expected utility of the heuristic.

The work described in this paper was inspired by a surprisingly effective neural network developed by Adorf and Johnston (1990) for scheduling astronomical observations on the Hubble Space Telescope. Our heuristic CSP method was distilled from an analysis of the network. In the process of carrying out the analysis, we discovered that the effectiveness of the network has little to do with its connectionist implementation. Furthermore, the ideas employed in the network can be implemented very efficiently within a symbolic CSP framework. The symbolic implementation is extremely simple. It also has the advantage that several different search strategies can be employed, although we have found that hill-climbing methods are particularly well-suited for the applications that we have investigated.

We begin the paper with a brief review of Adorf and Johnston’s neural network, and then describe our symbolic method for heuristic repair. Following this, we describe empirical results with the n -queens problem, graph-colorability problems and the Hubble Space Telescope scheduling application. Finally, we consider a theoretical model identifying general problem characteristics that influence the performance of the method. We include a second gratuitous citation to ourselves to illustrate a short citation (Minton et al., 1990).

66 2 Previous Work: The GDS Network

67 By almost any measure, the Hubble Space Telescope scheduling problem is
68 a complex task (Johnston, 1987; Waldrop, 1989). Between ten thousand and
69 thirty thousand astronomical observations per year must be scheduled, sub-
70 ject to a great variety of constraints including power restrictions, observation
71 priorities, time-dependent orbital characteristics, movement of astronomical
72 bodies, stray light sources, etc. Because the telescope is an extremely valu-
73 able resource with a limited lifetime, efficient scheduling is a critical concern.
74 An initial scheduling system, developed using traditional programming meth-
75 ods, highlighted the difficulty of the problem; it was estimated that it would
76 take over three weeks for the system to schedule one week of observations.
77 As described in section 4, this problem was remedied by the development
78 of a successful constraint-based system to augment the initial system. At
79 the heart of the constraint-based system is a neural network developed by
80 Johnston and Adorf (1989), the Guarded Discrete Stochastic (GDS) network,
81 which searches for a schedule.

82 From a computational point of view the network is interesting because
83 Adorf and Johnston found that it performs well on a variety of tasks, in
84 addition to the space telescope scheduling problem. For example, the network
85 performs significantly better on the n -queens problem than methods that
86 were previously developed. The n -queens problem requires placing n queens
87 on an $n \times n$ chessboard so that no two queens share a row, column or diagonal.
88 The network has been used to solve problems of up to 1024 queens, whereas
89 most heuristic backtracking methods encounter difficulties with problems
90 one-tenth that size (Stone and Stone, 1987).

91 The GDS network is a modified Hopfield network (Hopfield, 1982). In a
92 standard Hopfield network, all connections between neurons are symmetric.
93 In the GDS network, the main network is coupled asymmetrically to an
94 auxiliary network of *guard neurons* which restricts the configurations that
95 the network can assume. This modification enables the network to rapidly
96 find a solution for many problems, even when the network is simulated on

97 a serial machine. Unfortunately, convergence to a stable configuration is no
98 longer guaranteed. Thus the network can fall into a local minimum involving
99 a group of unstable states among which it will oscillate. In practice, however,
100 if the network fails to converge after some number of neuron state transitions,
101 it can simply be stopped and started over.

102 To illustrate the network architecture and updating scheme, let us con-
103 sider how the network is used to solve binary constraint satisfaction problems.
104 A problem consists of n variables, $X_1 \dots X_n$, with domains $D_1 \dots D_n$, and a
105 set of binary constraints. Each constraint $C_\alpha(X_j, X_k)$ is a subset of $D_j \times D_k$
106 specifying incompatible values for a pair of variables. The goal is to find
107 an assignment for each of the variables which satisfies the constraints. (In
108 this paper we only consider the task of finding a single solution, rather than
109 that of finding all solutions.) To solve a CSP using the network, each vari-
110 able is represented by a separate set of neurons, one neuron for each of the
111 variable's possible values. Each neuron is either "on" or "off", and in a solu-
112 tion state, every variable will have exactly one of its corresponding neurons
113 "on", representing the value of that variable. Constraints are represented by
114 inhibitory (i.e., negatively weighted) connections between the neurons. To
115 insure that every variable is assigned a value, there is a guard neuron for
116 each set of neurons representing a variable; if no neuron in the set is on, the
117 guard neuron will provide an excitatory input that is large enough to turn
118 one on. (Because of the way the connection weights are set up, it is unlikely
119 that the guard neuron will turn on more than one neuron.) The network is
120 updated on each cycle by randomly picking a set of neurons that represents
121 a variable, and flipping the state of the neuron in that set whose input is
122 *most inconsistent* with its current output (if any). When all neurons' states
123 are consistent with their input, a solution is achieved.

124 To solve the n -queens problem, for example, each of the $n \times n$ board posi-
125 tions is represented by a neuron whose output is either one or zero depending
126 on whether a queen is currently placed in that position or not. (Note that
127 this is a local representation rather than a distributed representation of the

board.) If two board positions are inconsistent, then an inhibiting connection exists between the corresponding two neurons. For example, all the neurons in a column will inhibit each other, representing the constraint that two queens cannot be in the same column. For each row, there is a guard neuron connected to each of the neurons in that row which gives the neurons in the row a large excitatory input, enough so that at least one neuron in the row will turn on. The guard neurons thus enforce the constraint that one queen in each row must be on. As described above, the network is updated on each cycle by randomly picking a row and flipping the state of the neuron in that row whose input is most inconsistent with its current output. A solution is realized when the output of every neuron is consistent with its input.

3 Why does the GDS Network Perform So Well?

Our analysis of the GDS network was motivated by the following question: “Why does the network perform so much better than traditional backtracking methods on certain tasks”? In particular, we were intrigued by the results on the n -queens problem, since this problem has received considerable attention from previous researchers. For n -queens, Adorf and Johnston found empirically that the network requires a linear number of transitions to converge. Since each transition requires linear time, the expected (empirical) time for the network to find a solution is $O(n^2)$. To check this behavior, Johnston and Adorf ran experiments with n as high as 1024, at which point memory limitations became a problem.¹

¹The network, which is programmed in Lisp, requires approximately 11 minutes to solve the 1024 queens problem on a TI Explorer II. For larger problems, memory becomes a limiting factor because the network requires approximately $O(n^2)$ space. (Although the number of connections is actually $O(n^3)$, some connections are computed dynamically rather than stored).

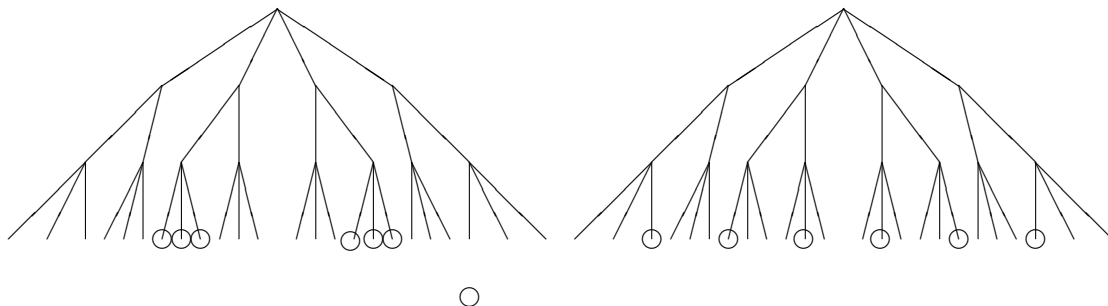


Figure 1: Solutions Clustered vs. Solutions Evenly Distributed

3.1 Nonsystematic Search Hypothesis

Initially, we hypothesized that the network's advantage came from the non-systematic nature of its search, as compared to the systematic organization inherent in depth-first backtracking. There are two potential problems associated with systematic depth-first search. First, the search space may be organized in such a way that poorer choices are explored first at each branch point. For instance, in the n -queens problem, depth-first search tends to find a solution more quickly when the first queen is placed in the center of the first row rather than in the corner; apparently this occurs because there are more solutions with the queen in the center than with the queen in the corner (Stone and Stone, 1987). Nevertheless, most naive algorithms tend to start in the corner simply because humans find it more natural to program that way. However, this fact by itself does not explain why nonsystematic search would work so well for n -queens. A backtracking program that randomly orders rows (and columns within rows) performs much better than the naive method, but still performs poorly relative to the GDS network.

The second potential problem with depth-first search is more significant and more subtle. As illustrated by figure 1, a depth-first search can be a disadvantage when solutions are not evenly distributed throughout the search space. In the tree at the left of the figure, the solutions are clustered together. In the tree on the right, the solutions are more evenly distributed.

172 Thus, the average distance between solutions is greater in the left tree. In a
173 depth-first search, the average time to find the first solution increases with the
174 average distance between solutions. Consequently depth-first search performs
175 relatively poorly in a tree where the solutions are clustered, such as that on
176 the left (Ginsberg and Harvey, 1990; Langley, 1992). In comparison, a search
177 strategy which examines the leaves of the tree in random order is unaffected
178 by solution clustering.

179 We investigated whether this phenomenon explained the relatively poor
180 performance of depth-first search on n -queens by experimenting with a ran-
181 domized search algorithm, called a Las Vegas algorithm (Brassard and Brat-
182 ley, 1988). The algorithm begins by selecting a path from the root to a leaf.
183 To select a path, the algorithm starts at the root node and chooses one of
184 its children with equal probability. This process continues recursively until a
185 leaf is encountered. If the leaf is a solution the algorithm terminates, if not,
186 it starts over again at the root and selects a path. The same path may be
187 examined more than once, since no memory is maintained between successive
188 trials.

189 The Las Vegas algorithm does, in fact, perform better than simple depth-
190 first search on n -queens (Brassard and Bratley, 1988). However, the perfor-
191 mance of the Las Vegas algorithm is still not nearly as good as that of the
192 GDS network, and so we concluded that the systematicity hypothesis alone
193 cannot explain the network's behavior.

194 3.2 Informedness Hypothesis

195 Our second hypothesis was that the network's search process uses informa-
196 tion about the current assignment that is not available to a constructive
197 backtracking program. 's use of an iterative improvement strategy guides
198 the search in a way that is not possible with a standard backtracking algo-
199 rithm. We now believe this hypothesis is correct, in that it explains why the
200 network works so well. In particular, the key to the network's performance
201 appears to be that state transitions are made so as to reduce the number of

202 outstanding inconsistencies in the network; specifically, each state transition
203 involves flipping the neuron whose output is most inconsistent with its cur-
204 rent input. From a constraint satisfaction perspective, it is as if the network
205 reassigns a value for a variable by choosing the value that violates the fewest
206 constraints. This idea is captured by the following heuristic:

207 **Min-Conflicts heuristic:**

208 *Given:* A set of variables, a set of binary constraints, and an assign-
209 ment specifying a value for each variable. Two variables *conflict* if
210 their values violate a constraint.

211 *Procedure:* Select a variable that is in conflict, and assign it a value
212 that minimizes the number of conflicts. (Break ties randomly.)

213 We have found that the network’s behavior can be approximated by a
214 symbolic system that uses the min-conflicts heuristic for hill climbing. The
215 hill-climbing system starts with an initial assignment generated in a prepro-
216 cessing phase. At each choice point, the heuristic chooses a variable that is
217 currently in conflict and reassigns its value, until a solution is found. The
218 system thus searches the space of possible assignments, favoring assignments
219 with fewer total conflicts. Of course, the hill-climbing system can become
220 “stuck” in a local maximum, in the same way that the network may become
221 “stuck” in a local minimum. In the next section we present empirical evi-
222 dence to support our claim that the min-conflicts approach can account for
223 the network’s effectiveness.

224 There are two aspects of the min-conflicts hill-climbing method that dis-
225 tinguish it from standard CSP algorithms. First, instead of incrementally
226 constructing a consistent partial assignment, the min-conflicts method *re-*
227 *pairs* a complete but inconsistent assignment by reducing inconsistencies.
228 Thus, it uses information about the current assignment to guide its search
229 that is not available to a standard backtracking algorithm. Second, the use
230 of a hill-climbing strategy rather than a backtracking strategy produces a
231 different style of search.

```

Procedure INFORMED-BACKTRACK (VARS-LEFT VARS-DONE)
  If all variables are consistent, then solution found, STOP.
  Let VAR = a variable in VARS-LEFT that is in conflict.
  Remove VAR from VARS-LEFT.
  Push VAR onto VARS-DONE.
  Let VALUES = list of possible values for VAR in ascending order according
                  to number of conflicts with variables in VARS-LEFT.
  For each VALUE in VALUES, until solution found:
    If VALUE does not conflict with any variable that is in VARS-DONE,
    then Assign VALUE to VAR.
      Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
    end if
  end for
end procedure

Begin program
  Let VARS-LEFT = list of all variables, each assigned an initial value.
  Let VARS-DONE = nil
  Call INFORMED-BACKTRACK(VARS-LEFT VARS-DONE)
End program

```

Figure 2: Informed Backtracking Using the Min-Conflicts Heuristic

232 3.2.1 Repair-Based Search Strategies

233 (This is a example of a third level section.) Extracting the method from the
234 network enables us to tease apart and experiment with its different compo-
235 nents. In particular, the idea of repairing an inconsistent assignment can be
236 used with a variety of different search strategies in addition to hill climbing.
237 For example, we can backtrack through the space of possible repairs, rather
238 than using a hill-climbing strategy, as follows. Given an initial assignment
239 generated in a preprocessing phase, we can employ the min-conflicts heuristic
240 to order the choice of variables and values to consider, as described in figure
241 2. Initially, the variables are all on a list of VARS-LEFT, and as they are
242 repaired, they are pushed onto a list of VARS-DONE. The algorithm attempts
243 to find a sequence of repairs, such that no variable is repaired more than
244 once. If there is no way to repair a variable in VARS-LEFT without violat-
245 ing a previously repaired variable (a variable in VARS-DONE), the algorithm
246 backtracks.

247 Notice that this algorithm is simply a standard backtracking algorithm
248 augmented with the min-conflicts heuristic to order its choice of which vari-
249 able and value to attend to. This illustrates an important point. The back-
250 tracking repair algorithm incrementally extends a consistent partial assign-
251 ment (i.e., VARS-DONE), as does a constructive backtracking program, but
252 in addition, uses information from the initial assignment (i.e., VARS-LEFT)
253 to bias its search. Thus, it is a type of *informed backtracking*. We still char-
254 acterize it as repair-based method since its search is guided by a complete,
255 inconsistent assignment.

256 4 Experimental Results

257 [section ommitted]

258 5 A Theoretical Model

259 [section ommitted]

260 6 Discussion

261 [section ommitted]

262 7 Acknowledgement

263 The authors wish to thank Hans-Martin Adorf, Don Rosenthal, Richard
264 Franier, Peter Cheeseman and Monte Zweben for their assistance and ad-
265 vice. We also thank Ron Musick and our anonymous reviewers for their
266 comments. The Space Telescope Science Institute is operated by the Associ-
267 ation of Universities for Research in Astronomy for NASA.

268 Appendix A. Probability Distributions for N- 269 Queens

270 [section ommitted]

271 References

- 272 Adorf, H. and Johnston, M. (1990). A discrete stochastic neural network
273 algorithm for constraint satisfaction problems. In *Proceedings of the Inter-*
274 *national Joint Conference on Neural Networks*.
- 275 Brassard, G. and Bratley, P. (1988). *Algorithmics - Theory and Practice*.
276 Englewood Cliffs, NJ: Prentice Hall.
- 277 Ginsberg, M. and Harvey, W. (1990). Iterative broadening. In *Proceedings of*
278 *AAAI-91*.

279 Hopfield, J. (1982). Neural networks and physical systems with emergent
280 collective computational abilities. In *Proceedings of the National Academy*
281 *of Sciences*, vol. 79, Washington, DC: National Academy Press.

282 Johnson, D., Papadimitrou, C. and Yannakakis, M. (1988). How easy is local
283 search? *Journal of Computer and System Sciences*, 37, pp. 79–100.

284 Johnston, M. (1987). Automated telescope scheduling. In *Proceedings of the*
285 *Symposium on Coordination of Observational Projects*.

286 Johnston, M. and Adorf, H. (1989). Learning in stochastic neural networks
287 for constraint satisfaction problems. In *Proceedings of NASA Conference*
288 *on Space Telerobotics*, vol. 37.

289 Langley, P. (1992). Systematic and nonsystematic search strategies. In *Pro-*
290 *ceedings of AAAI-92*.

291 Minton, S., Johnston, M., Philips, A. and Laird, P. (1990). Solving large scale
292 constraint satisfaction and scheduling problems using a heuristic repair
293 method. In *Proceedings of AAAI-90*.

294 Simmons, R. (1988). A theory of debugging plans and interpretations. In
295 *Proceedings of AAAI-88*.

296 Stone, H. and Stone, J. (1987). Efficient search techniques - an empirical
297 study of the n-queens problem. *IBM Journal of Research and Development*,
298 31, pp. 464–474.

299 Sussman, G. J. (1975). *A Computer Model of Skill Acquisition*. New York:
300 New American Elsevier.

301 Waldrop, M. (1989). Will the Hubble space telescope compute? *Science*, 243,
302 pp. 1437–1439.